

The trouble with quantifiers: Explaining children's deficits in scalar implicature

Alexandra C. Horowitz*

Department of Psychology, Stanford University

Rose M. Schneider*

Department of Psychology, Stanford University

Michael C. Frank

Department of Psychology, Stanford University

The two first authors contributed equally to this paper. We gratefully acknowledge Tamara Mekler for her assistance in data collection, the staff and families at Bing Nursery School, as well as David Barner for guidance on the Give-Quantifier task.

Address all correspondence to Rose M. Schneider, Stanford University, Department of Psychology, Jordan Hall, 450 Serra Mall (Bldg. 420), Stanford, CA, 94305. Phone: 650-721-9270. E-mail: rschneid@stanford.edu.

Abstract

Using language requires making pragmatic inferences that extend beyond the literal meaning of utterances. For example, “*some* of the cookies are oatmeal raisin” carries the scalar implicature that *not all* the cookies are. Adults make this pragmatic inference very reliably but children generally have more difficulty. Their performance varies tremendously across studies, however, and it is unclear how much of their difficulty stems from issues specific to quantifiers and how much comes from general pragmatic or processing demands. We designed an experimental paradigm to address this question. In Experiment 1, we measured children’s ability to compute both ad-hoc (contextual) and scalar (quantifier) implicatures. While 4-year-olds were at ceiling for ad-hoc descriptions, they performed poorly using quantifiers. We also found correlated performance between the quantifiers “some” and “none.” Experiment 2 replicated this correlation in an experiment including only quantifiers. In Experiment 3 we found that inhibitory control did not predict the ability to make implicatures, but that performance across quantifier tasks was highly correlated. Our findings suggest that difficulty with scalar implicatures may in fact be rooted in a lack of quantifier knowledge.

Keywords: scalar implicature, pragmatics, development

Introduction

In speech, adult listeners routinely make inferences that go beyond the literal sense of an utterance. For example, an adult who hears “I ate *some* of the cookies” would expect that the speaker did not eat *all* the cookies. Similarly, an adult listener who hears “I ate the sugar cookies,” would most likely assume that the speaker ate *only* the sugar cookies, and not the other varieties. These two statements use a weaker literal description (in these cases, scalar and contextual, respectively) to implicate that a stronger alternative is true. The first statement requires the listener to make a *scalar implicature*, which relies on lexical scales such as quantifiers (“some” vs. “all”) or modals (“possibly” vs. “definitely”) (? , ?). The second statement requires an *ad-hoc implicature*, in which a stronger description is negated by a contextually weaker description.¹

While adults sometimes incur a processing cost in computing scalar implicatures along quantifier lexical scales like <SOME, ALL> (? , ? , ? , ?), by and large they compute implicatures reliably in a wide variety of situations. In contrast, early investigations showed that children’s accuracy on these same scales was variable even until fairly late in development (? , ?). This result was intriguing because implicatures are an important case study for children’s pragmatic reasoning more generally. The finding that even elementary-aged children struggled with certain types of pragmatic judgments suggested a surprising disconnect between pragmatics and other aspects of language development.

While this early work made an important contribution by identifying an area of difficulty for children, both the use of a truth-judgment task and the abstract, propositional nature of the materials (e.g., “Some giraffes have long necks.”) may have lead the paradigm to underestimate

¹While ? (?) distinguished “generalized” and “particularized” implicatures, this distinction has been controversial. Here we use “ad-hoc implicature” as a term of convenience to describe contextually-supported inferences while remaining agnostic about the theoretical distinction.

children's ability to make implicatures. In more recent investigations, children have shown a graded pattern of successes and failures across different tasks (e.g., ?, ?, ?, ?). And other evidence has suggested early competence in making some pragmatic judgments. For example, five-year-olds show evidence of recognizing that pragmatically-infelicitous statements deserve smaller rewards, when given a graded—rather than binary—judgment task (?, ?). And some three- and four-year-olds are able to make ad-hoc (contextual) implicatures in a stripped down referent-selection paradigm (?, ?). So although some methods still show notable failures (e.g. ?, ?), the evidence is more mixed as to what children's abilities are and precisely what causes the observed deficits.

Although much progress has been made in this area, one weakness of the current literature is its diversity. Table 1 provides a (non-exhaustive) summary of a number of influential experimental papers. First, a wide range of methods and measures—from referent selection to felicity judgment—have been used to assess children's implicature abilities. Second, most previous studies have targeted relatively wide age bins (12–18 months or wider), ranges that do not allow for precise developmental comparisons between studies. And finally, other paradigm-level differences might lead to widely varying levels of performance. For example, recent work indicates that the use of the partitive (e.g., “*Some* of the cookies”) is a particularly strong cue for adults' scalar implicatures (?, ?), but studies vary in their use of this cue.

Study	Scale(s)	Ages	Measure	Sentence Type	Main Finding
Noveck (2001)	non-necessity–necessity, possibility–impossibility, some–all	5;1–5;11, 7;1–8;0, 9;0–9;5, 10;0–11;7	Truth Value Judgment (<i>Yes, I agree</i> or <i>No, I do not agree</i>)	“Some giraffes have long necks”	Comprehension matters: Children demonstrate pragmatic competence by age 7 in evaluating a variety of plausible and implausible sentences.
Papafragou & Mussolini (2003)	some–all, two–three, start–finish	4;11–5;11 (Study 1) 5;1–6;5 (Study 2)	Felicity Judgment (<i>Did Minnie answer well?</i>)	“Some of the horses jumped over the fence” (when all of the horses jumped over the fence)	Support matters: Children were more likely to reject infelicitous weak descriptions for numbers, and for all types of weak descriptions in the task with more pragmatic support (informativeness training, context of competition, statements about specific events).
Papafragou & Tantalou (2004)	some–all, ad-hoc, encyclopedic	4;1–6;1	Felicity Judgment (Decide whether or not to award a speaker a prize)	<i>Did you color the stars?</i> “I colored some” (when all were colored)	Scales matter: Children mainly withheld prizes for weak descriptions, and at higher rates for ad-hoc trials than other trial types.
Guasti, Chierchia, Crain, Foppolo, Gualmini, & Meroni (2005)	some–all	7;0–7;7	Truth Value Judgment: <i>Yes, I agree</i> or <i>No, I do not agree</i> (Studies 1–3), <i>Did Carolina say the wrong thing?</i> (Study 4)	“Some giraffes have long necks” (replication of Noveck (2001)), and scene descriptions, e.g. “Some monkeys are eating a biscuit” (when all are)	Context matters: 7-year-olds reliably computer implicatures for “some” after a training that increased their sensitivity to the informativeness of speakers’ descriptions (e.g., calling a grape “fruit” instead of “grape”) and when contextualized as evaluating a novice speaker novice speaker describing a scene.
Miller, Schmitt, Chang, & Munn (2005)	some–all	4;1–5;5 (Study 1) 3;6–5;10 (Study 2)	Direct Instruction Task (Study 1); Picture Matching Task (Study 2)	“Make some faces HAPPY/Make SOME faces happy/Make some HAPPY faces” (Study 1), “Shoe me where Pete made some faces HAPPY/Show me where Pete made SOME faces happy”	Prosody matters: In both tasks (completing the scene or selecting the referent), children reliably identified only a subset of the faces (out of four) when “some” was stressed, but not when it was unstressed.
Huang & Snedeker (2009)	some–all, two–three	5;2–6;1 (Study 1) 5;5–6;9 (Studies 2 & 3)	Eye-tracking referent selection	“Point to the girl with some of the socks” (when other girls and boys have shares of socks and soccer balls)	Time scale matters: Across studies, children were delayed in identifying the referent for scalar implicature trials, and accept and overlap between the meaning of “some” and “all.”
Katsos & Bishop (2011)	some–all, ad-hoc	5;1–6;3	Binary Truth Value Judgment (Study 1); Ternary Truth Value Judgment (Study 2); Sentence-to-picture Matching Task (Study 3)	“The mouse picked up some of the carrots”	Measures matter: While children tended to accept under-informative scalar and ad-hoc descriptions given a binary decision, they showed sensitivity to weaker statements given a ternary choice or picture matching task.
Barner, Brooks, & Bale (2011)	some–all, ad-hoc	4;0–5;0	Truth Value Judgment	“Are some of the animals sleeping?” (when all are)	Specificity matters: 4-year-olds accept weak ad-hoc and scalar descriptions. When preceded by restrictive “only”, they reject ad-hoc descriptions but continue to accept that “only some” can mean <i>all</i> .
Skordos & Papafragou (2014)	some–all	4;9–5;8	Felicity Judgment (<i>Did the puppet answer well?</i>)	“Some of the blickets have a crayon” (when all of them do)	Comparisons matter: Children were more likely to reject infelicitous uses of “so” if they first heard “all” falsely refer to quantity (only 3/4 blackouts had crayons), but not is “all” referred falsely to the objects (e.g., “all of the blackest have a scarf”).

Table 1
Review of previous literature on children’s comprehension of implicatures.

Thus, one goal of our current work is to consolidate insights from the previous literature by providing a direct comparison between ad-hoc and scalar implicatures in a simple referent selection task. This comparison allows for strong inferences about developmental differences between making scalar and ad-hoc implicatures. Additionally, a second goal of our current work is to explore theoretical hypotheses about the source of children's surprising failures in scalar implicatures particularly. (We follow previous work in suggesting that, if children succeed in ad-hoc implicatures and fail in scalar implicatures, it is very unlikely that a general pragmatic deficit explains previous findings.) The next section describes possible candidate explanations.

Candidate Explanations for Failures in Scalar Implicature

One candidate explanation for children's failures to compute scalar implicature has emerged from previous literature. This idea, known as the *Alternatives Hypothesis*, suggests that their children's performance may have less to do with general pragmatic knowledge per se, and more to do with their knowledge of the particular scales on which implicatures are computed. Barner and colleagues (2006, 2007, 2008) suggested that children's ability to compute scalar implicatures relies on their recognition of the relevant lexical alternatives (e.g., that use of the weaker term "some" conveys a direct contrast with the stronger alternative "all", thus implying *some but not all*). In other words, children's pragmatic inferences rely on their ability to consider relevant possible alternative word choices that could have been used in place of the ones the speaker chose. So even in supportive paradigms, if children cannot bring to mind "all" when reasoning about "some," they will fail to make an implicature.

Previous research has provided a number of tests of this hypothesis. First, children's performance in implicature tasks increases when they have stronger access to specifically lexical alternatives. ? (2008) made scalar alternatives salient within the experimental paradigm via a blocked

design (e.g. “all” trials presented before “some” trials) and found that performance increased. Second, even when lexical alternatives are not presented directly, the use of more supportive paradigms or designs could be helpful. In particular, if referents corresponding to particular lexical alternatives are present on each trial, the presence of these referents could facilitate reasoning about alternatives. Supporting this idea, preschoolers showed some preliminary evidence of computing scalar implicatures for quantifiers in a forced-choice paradigm where different quantification scenarios were pictured (?, ?). And children were able to reason about ad-hoc scales both in a forced-choice paradigm (?, ?) and in a truth-value paradigm (?, ?) when the relevant alternative objects were present in the trial context.

On the other hand, this body of research leaves open an alternative explanation. Perhaps the problem is not with inferential alternatives generally, but instead a specific issue with quantificational alternatives. Perhaps children have trouble summoning to mind quantificational alternatives because at varying developmental time points they either 1) do not know the quantifiers, 2) do know them but cannot process them effectively, or 3) know them and can process them but have not grouped them into sets of lexical alternatives. In some sense, all three of these possibilities are consistent with the general spirit of the Alternatives Hypothesis, but problems 1 and 2 are perhaps more specific to quantifier meaning than otherwise supposed.

Supporting this general line of reasoning, quantifiers are a difficult lexical class for children to acquire. Unlike many other lexical classes (e.g., colors), quantifiers are not mutually exclusive with one another. In addition, quantifier semantics are relative to the total set size. For example, in a set size of 8, the quantifier “some” can be used felicitously to describe anywhere from 2 – 7 objects (?, ?) (and truth-functionally to describe sets 1 – 8), although adult-like responses tend to center around the mean of the set-size (?, ?). These difficulties can be seen in an experiment by ?

(?), who found that 3–4-year-olds were only 75% correct in choosing the referent of “all” in a four-alternative forced choice. Thus, perhaps quantificational alternatives in particular general are the problem, rather than alternatives more generally.

The Current Study

In three experiments, we explore these ideas about the determinants of performance in implicature tasks using a novel paradigm. We designed a simple referent selection task in which children were asked to select which of three book covers they thought the experimenter was describing. This task gave children access to the visual alternatives in each trial (the three book covers) and the lexical alternatives across trials (either ad-hoc or scalar descriptions). Our design allowed us to counterbalance trial types (ad-hoc vs. scalar descriptions crossed with implicature vs. unambiguous control targets) fully across participants, to examine both within-subject patterns of responses and between-subject developmental patterns, while also reducing the demands of the task.

In Experiment 1, we included both ad-hoc and scalar descriptions with implicature and control trials for each. Four-year-olds were at ceiling on ad-hoc trials (similar to previous work, e.g. ?, ?), but their performance on scalar implicature trials was very low. In Experiment 2, we omitted ad-hoc trials. We found developmental increases in performance for each trial type, with higher performance on implicature trials for 4-year-olds in this scalar-only version. In both Experiments 1 and 2, we found an unexpected result: Children’s pattern of responses on scalar implicature trials was bimodal and strongly correlated with their performance on “none” (scalar control) trials. This correlation suggested a general source of difficulty in comprehending these quantifiers beyond simply failing to make a scalar implicature.

In Experiment 3, we used an individual differences study to explore this correlation further.

To ask whether our measurement of quantifier knowledge was related to the specific task we used, we measured quantifier knowledge with a separate paradigm (?, ?). We additionally assessed the alternative hypothesis that inhibitory control development, rather than quantifier alternative knowledge, might underlie the correlation between “none” and implicature performance; we found no evidence for this interpretation. Overall, our findings suggest that while preschoolers’ computation of scalar implicatures can be supported by stronger recognition of the lexical alternatives, their failures may be rooted in difficulty comprehending and contrasting relevant quantifiers.

Experiment 1: Ad-hoc and scalar implicature computation in children

Given the difficulty equating results on children’s computation of implicatures across different methods and paradigms, we created a single task that could be adapted to investigate both ad-hoc and scalar items in one task. This task involved one set of visual stimuli presented in the same order to all participants; however, the particular items (ad-hoc or scalar) queried were counterbalanced across participants. Thus, with one set of visual stimuli we could directly compare children’s performance on both ad-hoc and scalar implicatures in a single experimental session. In Experiment 1, we included both ad-hoc and scalar implicatures. In all scalar items, we used a partitive construction to increase the likelihood of making an implicature (?, ?).²

Methods

Participants. Table 2 shows the age information for a planned sample of 48 children recruited from a university preschool. Two children were excluded from the final sample for not

²All stimuli, data, and analyses are available in a version-controlled public repository at http://github.com/rosemshneider/si_paper.

Age group	N	Mean	Median	SD
4.0 – 4.5-year-olds	24	4.2	4.19	0.14
4.5 – 5.0-year-olds	24	4.74	4.73	0.16

Table 2
Participant age information for Experiment 1.

completing the task, and one additional child was excluded due to experimenter error. All participants' primary language was English, and no child completed more than one session of the task.

Stimuli. Experimental stimuli consisted of 18 sets of three printed pictures of book covers, each featuring four familiar items. In each trial, one book cover contained four items of the same kind (e.g., four cats), another book cover contained four items of another kind (e.g., dogs), and the final cover contained two items of a new set and two items repeated from one of the other book covers (e.g., two birds and two cats). An example of the stimuli can be seen in Figure 1. All items on the book covers were familiar to children, and were able to be identified. All participants saw the book covers in the same order.

Procedure. Participants were tested in individual sessions in a quiet room at their nursery school. The experimenter introduced the study as a guessing game, and explained that the child would receive a hint about which book cover the experimenter had in mind. The experimenter emphasized that the child would only receive one clue about what book the experimenter was describing, and she had to use that clue to make her decision. All participants saw image sets (three books) in the same order; however, these image sets were counterbalanced for target location across the three scripts. Description condition and trial-type were further randomized




Figure 1. Example trial stimuli used in all experiments. Children received a clue from the experimenter about which book she had in mind and responded based solely on the clue; this was either an ad-hoc or a scalar description of a book with either an unambiguous or implicature target.

across participants, and were spaced to avoid immediate repeat trial types. Table 3 shows the breakdown of trial types and sample scripts. Children did not receive feedback after the test trial.

Prior to the test trials, children were familiarized to the task with a practice trial with three book covers, each displaying a single unique and familiar item. At the start of the practice trial, the experimenter told the child “On the cover of my book, there’s a TV.” After children successfully completed the practice, they saw 18 test trials. At the beginning of every trial, the experimenter provided the child with either an ad-hoc or scalar description of one book, and instructed the child to point to the book she was describing. If the child pointed to more than one book, or the response was otherwise ambiguous, the experimenter emphasized again that she was talking about just one book, and that the child should choose the single book she was describing.

In ad-hoc trials (eight total), the experimenter’s descriptions of the target book used the

Condition	Trial type	# trials, Expt. 1	# trials, Expts. 2 & 3	Statement: “On the cover of my book, ...”	Target
Scalar	implicature	4	6	“...some of the pictures are cats”	B
	all	2	6	“...all of the pictures are cats”	C
	none	2	6	“...none of the pictures are cats”	A
	unambiguous ‘some’	2		“...some of the pictures are birds”	B
Adhoc	implicature	4		“...there are cats”	C
	distractor	2		“...there are dogs”	A
	comparison	2		“...there are birds”	B

Table 3

Study designs for Experiments 1, 2, and 3, using script examples for the stimulus set pictured in Figure 1.

names of the pictured objects, providing contextual support for the target. Ad-hoc control trials referred to an unambiguous target (e.g., “On the cover of my book, there are dogs” in Figure 1), while implicature trials required the child to reason about the speaker’s meaning given an ambiguous utterance (e.g., “On the cover of my book, there are cats,” which could refer to either the book containing only cats or the book containing cats and birds). In these critical trials, children had to understand that the speaker could potentially be talking about either the book with four or two of the named object, but that by opting to describe only one kind of object she was referring to the cover with four of the same object; otherwise, she would have mentioned both kinds of objects, or the ones unique to that cover (i.e., birds).

In scalar trials (ten total), the experimenter described the target book with quantifiers. For scalar items, control trials referred to unambiguous targets with the quantifiers “all” and “none” (e.g., “On the cover of my book, *all/none* of the pictures are cats”) or an unambiguous referent of “some” (e.g., “On the cover of my book, *some* of the pictures are birds.”). On critical scalar




Figure 2. Proportion of correct responses by each age group across all trial types and split by implicature type. Error bars show 95% confidence intervals computed by non-parametric bootstrap.

implicature trials, the experimenter used the weak quantifier “some” to reference the item pictured across two book covers (e.g., “On the cover of my book, *some* of the pictures are cats.”). These trials required the child to reason that because the speaker used the weak quantifier “some,” she must be referring to the book picturing only two of the named target, or else she would have used the stronger quantifier “all.”

Results

We found that all children performed at ceiling on ad-hoc implicature trials. In contrast to this success, however, they struggled significantly on implicature (“some”) and “none” trials. Children’s accuracy on all trial types is plotted in Figure 2. On implicature trials, children’s performance was coded as correct if they selected the image consistent with the implicature: the single item (e.g., cats) on the ad-hoc trials, and the mixed item (e.g., cats and birds) on the scalar trials. Children were at ceiling making ad-hoc implicatures, which is consistent with previous research suggesting that children are able to succeed making such implicatures when they have access to the relevant lexical alternatives (?, ?). Children’s performance across ad-hoc trials provides strong evidence that our novel paradigm is an appropriate measure for such items. In contrast to their success in making ad-hoc implicatures, however, children struggled on quantifier trials, performing much worse for both “some” and “none” trials.

We ran a logistic mixed effects model, predicting a correct response as an interaction of age, condition (ad-hoc or scalar) and trial type (implicature or control), with random effects of participant and trial type. We found that performance was slightly lower for scalar trials than ad-hoc trials ($\beta = -8.02, p = .09$), and that there was a significant interaction between condition and trial type, such that performance was significantly worse on scalar implicature trials ($\beta = 16.07, p < .02$). We also found a significant 3-way interaction between condition, trial type, and age, such that performance on scalar implicature trials decreased with age ($\beta = -4.16, p < .01$). There were no significant effects of adding trial order (trials in the first half vs. second half of the experiment), indicating that performance did not change throughout the course of the experiment.

In post-hoc analysis of the data, we found an unpredicted consistency in performance on




Figure 3. Frequency of mean performance, split by condition.

“some” and “none” trials (Figure 3).³ To examine this patterns of responses more closely, we ran Hartigan’s dip test and found significant bimodal distributions for both “some” ($D = .15, p < .0001$) and “none” ($D = .20, p < .0001$). This result suggests children did not respond at chance in scalar trials, but instead were consistently either correct or incorrect. Additionally, children’s success on “some” and “none” trials was correlated ($r = .45, p < .001$), such that children who performed

³We also observed this bimodal performance in a pilot sample ($N = 23$), so we had some reason to expect it in Experiment 1.

better on “some” trials also tended to perform better on “none” trials. Performance on “none” and “all” trials ($r = .11, p = .45$) and “some” and “all” trials ($r = .01, p = .95$) was not correlated.

Discussion

The results of Experiment 1 indicated that children were easily able to make ad-hoc, but not scalar, implicatures in our task. This pattern of performance was puzzling, given both our efforts to reduce task demands and children’s striking success in ad-hoc trials. Despite having access to both visual and lexical alternatives across the task, children still struggled to make scalar implicatures. In addition, performance did not increase across the course of the study, suggesting that the repetition of scalar alternatives across trials did not lead to greater levels of performance.

Even more intriguing was the unexpected developmental change we observed on “none” trials. We included “none” as an unambiguous control quantifier, but found that children had lower performance for this scalar term as well. These results are supported by previous work suggesting that even older preschoolers struggle with negation occurring in contexts without pragmatic support (?, ?).

It is possible that this pattern of performance stems from including both ad-hoc and scalar quantifier descriptions within one experimental session. Children’s success on ad-hoc trials may have lead to a misinterpretation of scalar descriptions (e.g., “On the cover of my book, some of the pictures are cats”) to ad-hoc descriptions (“On the cover of my book, there are cats”). Presenting scalar descriptors in the same experimental session as other relevant and felicitous descriptions may alter scalar implicature comprehension, even in adults (cf. ?, ?). To explore the whether children’s performance on scalar trials was influenced by ad-hoc trials, we removed all ad-hoc trials from our task, and ran a scalar-only version of the the study.

Age group	N	Mean	Median	SD
3.0 – 3.5-year-olds	12	3.4	3.35	0.1
3.5 – 4.0-year-olds	12	3.8	3.67	0.12
4.0 – 4.5-year-olds	14	4.3	4.24	0.1
4.5 – 5.0-year-olds	12	4.8	4.63	0.15

Table 4
Participant age information for Experiment 2.

Experiment 2

In Experiment 2, we pursued the possibility that children failed to make scalar implicatures as a result of competing ad-hoc descriptors in the same experimental session. Additionally, we expanded our target age-range to 3–5 years to more fully explore the developmental trajectory associated with making scalar implicatures.

Methods

Participants. Table 4 shows the age information for a new sample of 50 participants from a university preschool. One additional child was run but excluded for stopping the task early.

Stimuli. Stimuli were identical to Experiment 1. The only changes made in experimental protocol were to the scripts; all 18 test trials were converted to quantifier descriptions (Table 3). In Experiment 2, the 18 test trials consisted of six control “all” trials (e.g., “On the cover of my book, *all* of the pictures are cats”), six “none” trials (e.g., “...*none* of the pictures are cats”), and six “some” (scalar implicature) trials (“...*some* of the pictures are cats”). We removed the unambiguous “some” trials to more effectively counterbalance; in “some” trials, the quantifier always referenced the item pictured across two book covers (e.g., in Figure 1, children heard

references to “none,” “some,” or “all” cats). As in Experiment 1, image sets were presented in a fixed order, counterbalanced for both target location and book triad order. Participants were randomly assigned to one of three scripts, with a pseudo-randomized trial order such that every book set was referred to by each quantifier type, and the same trial type never immediately repeated. These three scripts were counterbalanced across participants.

Procedure. The procedure was identical to Experiment 1.

Results

In Experiment 2, children’s performance in all trial types increased with age (Figure 4). Performance was highest in “all” trials across all age groups. However, performance was still significantly lower in both “none” and “some” trials (in comparison to “all,” $p < .05$ for all tests); only 4.5–5-year-olds’ performance for “none” was not significantly different than for “all” ($t(11) = 1.74, p = .11$). Children’s performance in Experiment 2 was numerically, but not significantly different than in Experiment 1 in independent sample t-tests by trial type between age groups in both experiments ($p > .09$ for all tests). All of these tests were relatively low in power, however, due to the small number of individuals in each bin.

To aggregate across groups, we ran a planned logistic mixed effects model, predicting correct responses as an interaction of age and trial type (*all*, *some*, or *none*), with random effects of trial type by participant. The only significant effect that emerged was age, such that performance increased across trials as children got older ($\beta = 20, p < .001$). Adding trial order (first or second half of the experiment) to the model did not interact with any of the variables, indicating the performance did not change over the course of the experiment. We suspected that this lack of a main effect was due to individual variability, such as in Experiment 1.

Consistent with the findings from Experiment 1, we ran Hartigan’s dip test and again found




Figure 4. Proportion of correct responses by each age group across all scalar trial types. Error bars show 95% confidence intervals computed by non-parametric bootstrap.

significant bimodal patterns of responses for both “some” ($D = .12, p < .0001$) and “none” ($D = .15, p < .0001$) trials. Once again, these trial types were highly correlated with one another ($r = .52, p < .001$).

Thus, as an exploratory analysis, we ran another version of the mixed effects model removing the random effect of trial type, as we hypothesized that our initial model did not find trial type effects due to the correlational structure observed between trial types ($\text{correct} \sim$




Figure 5. Children’s selections on incorrect trials, with panel showing different trial types (“all,” “some,” and “none”). Bars represent the count of children choosing each incorrect alternative.

trial type * age + (1 | subject)). In addition to a main effect of age ($t = 1.88, p < .01$), this model revealed a conditional effect: “some” trials were lower than “all” trials ($t = -7.69, p < .01$), and marginally reduced from “none” trials ($t = 3.03, p = .09$). It also showed interactions between trial type and age, such that there was a greater difference between younger children’s performance on “some” and “all” trials ($t = 2.84, p < .001$), and “some” and “none” trials ($t = 0.90, p = .05$). Overall, we observed large individual variability in children’s performance, with

mean trends of children struggling with the quantifiers “some” and “none” in relation to “all.”

In a further exploratory analysis, we investigated the particular kinds of errors that children made in “some” and “none” trials (Figure 5). We found that children selected the correct noun on both kinds of trials, and chose the “all” option most frequently.

Discussion

In Experiment 1, we observed success in children’s computation of ad-hoc, but not scalar implicatures. To explore whether children’s performance in making scalar implicatures was hindered by the presence of both ad-hoc and scalar items in the same session, we excluded ad-hoc items. When presented with only scalar descriptions in Experiment 2, children’s performance was numerically (although not significantly) better than in Experiment 1, with scalar implicature performance positively correlated with age. We still observed low performance in “some” and “none” trials. Additionally, we again found bimodal and correlated patterns of responses in these two trials, with children consistently failing or succeeding in both “some” and “none” trials. In sum, the results of Experiment 2 indicated that children struggle with making scalar implicatures beyond dealing with competing contextual descriptors. The cause of this failure is not clear however, and individuals differed substantially.

Given children’s difficulty with “none” and the strong positive correlation between “some” and “none” trials, it is possible that making implicatures necessitates some familiarity with *both* ends of the quantifier scale (*none – some – all*). This idea is not consistent with classic pragmatic theory (e.g., ?, ?), which posits that only alternatives that logically entail the current quantifier (e.g. “all” or “most”) take part in the implicature computation. Nevertheless, some recent work supports this possibility: ? (?) found in a model of pragmatic felicity that “none” was heavily weighted as an alternative in the scalar pragmatic computation for “some.” So there is some

indirect evidence that children might need to know “none” to be able to make a scalar implicature with “some.”

In addition to understanding the extremes of the quantifier scale, another other possibility might also account for children’s failures. Children hearing “none of the pictures are cats” might simply match the word “cats” to the referent with the cats and fail to inhibit this match in favor of the correct alternative. This possibility is consistent with some work with adults on the comprehension of negation, where comprehenders have been posited to generate the positive match and then negate it (e.g. ?, ?). We explored these two alternatives in Experiment 3, including measures of both inhibitory control and quantifier knowledge in the same session.

Experiment 3: Inhibitory control and quantifier knowledge measures

In an individual differences paradigm, we supplemented our implicature task with two additional tasks: an inhibitory control task, the Dimensional Change Card Sort (DCCS) (?, ?); and a quantifier-knowledge task, Give-Quantifier (?, ?). The DCCS is a standard executive function measure that requires children to shift tasks midway through the task (e.g., sorting cards based on shape rather than color). Children’s performance in the DCCS (i.e., their ability to task switch) is a reliable measure of their inhibitory control (?, ?). In the Give-Quantifier task, a productive measure of quantifier knowledge, children give a quantity of items in response to a quantifier prompt. Thus, with these two tasks, we can assess the contributions of both quantifier knowledge and inhibitory control in driving children’s observed performance in our scalar implicature task.

Methods

Participants. Table 5 shows the age information for a new planned sample of 72 children from a university preschool; this sample was selected to have 80% power to detect correlations of

Age group	N	Mean	Median	SD
3.0 – 3.5-year-olds	18	3.23	3.24	0.13
3.5 – 4.0-year-olds	18	3.73	3.8	0.18
4.0 – 4.5-year-olds	18	4.16	4.11	0.11
4.5 – 5.0-year-olds	18	4.7	4.68	0.14

Table 5
Participant age information for Experiment 3.

$r > .3$. Twelve children additional were recruited but excluded from the final sample for having participated in either Experiments 1 or 2. Nine children asked for a break, and completed one of the three tasks in a subsequent testing session; these children were not excluded from analyses.

Stimuli. Stimuli for the implicature task were identical to Experiment 2. The materials for the DCCS, our inhibitory control measure, were drawn from the original methods paper (?, ?). Fourteen laminated sorting cards (7 red rabbits and 7 blue boats) were put into two plastic sorting trays marked with either a target blue rabbit or red boat. To assess quantifier knowledge, we used the Give-Quantifier task (?, ?). Stimuli for this task consisted of three different sets of plastic fruits (8 oranges, 8 bananas, and 8 strawberries) and a red plastic plate. Fruits were grouped together by kind at the start of each trial.

Procedure. Task order was counterbalanced across participants, and individual scripts for each task were also counterbalanced to avoid order effects. The tasks were done in a small room apart from the main classroom in individual sessions. The procedure for our implicature task was identical to Experiment 2. The experimenter asked the child before the start of every task whether she would like to play the game or return to the classroom.

We drew our protocol for DCCS directly from the original methods paper (?, ?). Children were shown two plastic trays each marked with a target card (a blue rabbit and a red boat). At the beginning of the task, the experimenter explained that this was either the shape or color game, and that the cards had to be sorted according (e.g., in the color game, a red rabbit would be sorted into the red boat tray). After six trials, the experimenter told the participant that the rules had changed, and the cards had to be sorted by the other dimension (e.g., after the switch to the shape game, a red rabbit would be sorted into the blue rabbit tray).

In running the Give-Quantifier task, we followed the protocol of the original study (?, ?), with the exception of limiting the quantifiers used in the task to “some,” “all,” “none,” and “most.” The experimenter used the partitive construction and prosodically emphasized the quantifier across all trials (e.g., “Can you put *all* of the bananas into the plate?”). Quantifiers were presented in two different orders between participants, and fruit-quantifier pairings were quasi-randomized such that the same pairing was not repeated within a session. If the child requested clarification, the experimenter repeated the prompt, and added that the child should put however many pieces of fruit she felt should go on the plate.

In coding the results of the Give-Quantifier task, we relied on the original coding scheme (?, ?); “all” and “none” trials were coded as correct for 8 and 0 pieces of fruit given respectively; “some” trials were coded as correct if the child gave between 2 and 7 pieces, and “most” trials were correct if the child gave between 5 and 7 pieces.

Results

We again replicated children’s performance on our implicature task (Figure 6). We found significantly higher performance for all age groups with the quantifier “all” versus “some” and “none” in independent sample t-tests ($p < .02$ for all tests), and again found developmental




Figure 6. Proportion of correct responses by each age group across all trial types for the scalar implicature task. Error bars show 95% confidence intervals computed by non-parametric bootstrap.

increases in performance for each quantifier. Experiment 2 was not different than Experiment 3: Independent sample t-tests between age-groups for Experiments 2 and 3 did not yield any significant differences ($p > .1$ for all tests) except for 4–4.5-year-olds’ performance on “all” trials, which was significantly lower in Experiment 3 ($t(30) = 2.16, p < .05$).

As in Experiment 2, we ran Hartigan’s dip test in a post-hoc analysis and again found a significant bimodal distribution of performance in both “some” ($D = .07, p = .002$) and “none”

trials ($D = .15, p > .0001$). Performance with these two quantifiers was also significantly positively correlated ($r = .4, p < .001$). Because children's performance was so highly correlated with age, we ran a partial correlation controlling for age, and found these trial types were still significantly correlated ($r = .3, p < .01$).

We next turned to whether children's lower and correlated performance on "some" and "none" trials was due to an inhibitory control issue (i.e., making a response based solely on the target noun, regardless of quantifier used.) Overall, we found a developmental increase in performance from three to five years of age. Performance on post-switch trials was significantly correlated with age ($r = .28, p = .018$), with 3–3.5-year-olds at chance (t (only 4–5-year-olds performing significantly better than chance (4–4.5-year-olds: $t(18) = 3.05, p < .01$; 4.5–5-year-olds: $t(16) = 3.31, p < .01$). After controlling for age, we did not find a significant correlation between inhibitory control and performance on either "some" ($r = .13, p = .26$) or "none" trials ($r = -.01, p = .93$) in our implicature task.

Next, we turned our attention the relationship between children's performance on our scalar implicature task and quantifier knowledge. Children's performance on the Give-Quantifier task was very similar to performance on our implicature task, with all age groups performing at ceiling for the quantifier "all", and only older children succeeding on "most," "none," and "some" quantifier trials. Figure 9 shows the breakdown of how children responded to these scalar terms. We collapsed across all age groups and found a significant bimodal distribution of responses through Hartigan's dip test in both "some" ($D = .19, p < .0001$) and "none" trials ($D = .15, p < .0001$). In an exploratory analysis, we ran dip tests by age group, and found that this effect was primarily driven in "some" trials by 3–3.5-year-olds ($D = .14, p < .0004$) and 4–4.5-year-olds ($D = .13, p < .004$), and in "none" trials by 3–3.5-year-olds ($D = .21, p < .0001$) and 3.5–4-year-olds




Figure 7. Proportion of correct responses by each age group for the Give Quantifier task, with quantifier prompts *all*, *most*, *none* and *some*. Error bars show 95% confidence intervals computed by non-parametric bootstrap.

($D = .2, p < .0001$).

Overall, we found that younger children showed a bimodal and correlated pattern of response to the quantifier “none,” with the majority of children giving either 0 or all 8 items in these trials, and gradually shifting to the correct response by 4.5–5 years. Similar to performance on “some” trials in our scalar implicature task, we found that younger children gave all 8 objects in response to the prompt “some,” and only the oldest age groups gave a more adult-like

distribution of items in response. In a partial correlation controlling for age, we found that performance in “none” and “some” trials was significantly correlated ($r = .61, p < .0001$).

In partial-correlations controlling for age, we found that performance on Give-Quantifier and implicature tasks was correlated. Children who tended to struggle with scalar terms in the context of implicatures also had lower performance when asked to produce a number of items in response to a quantifier prompt, specifically on “some” ($r = .27, p < .02$) and “none” trials in both tasks ($r = .52, p < .0001$). We did not find a significant correlation with performance on “some” scalar implicature and “none” Give-Quantifier trials ($r = .18, p = .14$), although we did find a relation between performance on “none” scalar implicature and “some” Give-Quantifier trials ($r = .35, p = .002$).

In a further exploration of the relationship between quantifier knowledge and scalar implicature performance, we examined both the particular kinds of errors that children made across both tasks, and how they were related. Figure 8 shows the breakdown of children’s performance in our scalar implicature task on incorrect trials. As in Experiment 2, children exhibited evidence of making selections based solely on the target noun, regardless of the quantifier used. This result closely mirrors children’s performance in our Give-Quantifier task (Figure 9), in which younger children responded in a bimodal fashion for the items “some” and “none.”

Discussion

In Experiment 3, we explored potential factors behind children’s difficulty making scalar implicatures, as well as with comprehension of the quantifier “none.” We combined two tasks targeting specific hypothesized areas of difficulty (inhibitory control and lack of quantifier knowledge) with our implicature paradigm in a within-subjects design to explore the relationship




Figure 8. Scalar implicature error analysis. Count of children choosing an alternative on incorrect trials, faceted by trial type (“all,” “some,” and “none”) and split by age group.

amongst these abilities. While we found that younger children did have difficulties in our inhibitory control task, we did not find a significant relation between performance in this task and inhibitory control, at least at the level that we had statistical power to detect. We did find that children’s performance on the Give-Quantifier and scalar implicature tasks were related, however, and children who failed “some” and “none” trials tended to do so across both tasks. This result indicates that quantifiers may be particularly difficult for children, even when removed from the




Figure 9. Give-Quantifier performance. Proportion of children giving numbers of items faceted by age group and quantifier prompts (“all,” “none,” and “some”).

particular task we designed.

While it is clear that children struggle with quantifier comprehension, the source of this developmental difficulty is less clear. Quantifiers are a difficult class of linguistic expression to acquire, due to their non-exact and varying corresponding set sizes (?, ?). Because children’s definitions of these scalar terms are not solidified in the pre-school years, it is possible that that they are not able to use and contrast quantifiers on the <NONE – SOME – ALL> scale to succeed

in make a scalar implicature.

General Discussion

We designed a simple task to investigate patterns of pragmatic development both within- and between-subjects. We minimized task demands by asking participants to select the speaker's intended referent from among three visual alternatives. In Experiment 1, we replicated the finding that preschoolers can compute ad-hoc implicatures (?, ?), though we found poor performance on scalar implicatures. In Experiment 2, we removed competing ad-hoc descriptions from our implicature task, and found marginal increases in preschoolers' comprehension of all scalar quantifier terms. We still observed correlated difficulty with the quantifiers "some" and "none," however. In Experiment 3, we explored two possible explanations for this pattern of performance, inhibitory control and quantifier knowledge, and found evidence that children's difficulties on our implicature task are rooted in a lack of quantifier knowledge, rather than inhibitory control. Our findings suggest that while 4-year-olds are able to compute scalar implicatures with support from contextual cues, their performance is fragile.

Our work contributes to the existing literature in a number of ways. First, it offers a novel paradigm that is less complicated than many other implicature tasks, leading us to feel more confident that our results reflect children's true sensitivity rather than inadvertent task demands. Each test set remained visible to children, and they were merely asked to select which picture they thought was the referent corresponding to the speaker's description. Second, the relatively high number of trials for each participant both helped increase the precision of our measurements and also offered the possibility for children to identify lexical alternatives as the study progressed. Third, we were able not only to compare performance across age groups, but also to examine individual patterns of responses across the different trial types. This design helped us determine

that preschoolers' performance on scalar implicature trials was bimodal and highly related to their performance on "none" trials, which we would have been unlikely to uncover in a purely between-subjects implicature design without controls. Finally, we were able to test two hypotheses about the sources of children's difficulty with scalar implicatures, and found evidence for a link between poor quantifier knowledge and the ability to make scalar implicatures.

Our findings provide some support for the Alternatives Hypothesis (?, ?, ?). In particular, our ad-hoc trials in Experiment 1 show that preschoolers had no difficulty generally making inferences about contextual descriptions when alternative nominal descriptions were obvious from the context. In addition, the presence of ad-hoc trials decreased scalar performance in Experiment 1 compared with Experiment 2, suggesting that children might have recognized the ad-hoc descriptions as alternatives that competed with the quantifiers.

On the other hand, another pattern in our data was more difficult to reconcile with the Alternatives Hypothesis: Children's performance did not change over the course of either experiment. We had expected that, if children's difficulties with scalar implicature were due to a lack of recognition of the contrastive relationship between "some" and "all," that this relationship would be revealed by the two words' consistent use in contrasting references over the course of the many trials that each child completed (cf. ?, ?). The lack of trial order effects we observed could indicate that children in our task did not yet have strong enough comprehension of these terms for contrastive use to matter, or alternatively that our referent-selection task eliminated the problem of summoning the contrasting term to mind and instead foregrounded other inferential issues.

Further, the correlated responses for "some" and "none" trials in both Experiments 1 and 2 suggested that children's difficulty with these quantifiers could have a common root, either in quantifier knowledge or inhibitory control. Our data in Experiment 3 provided evidence against a

purely inhibitory account, leaving the quantifier knowledge hypothesis as one possible contender. Although “none” is not typically considered part of the same Horn scale as “some” and “all,” it is nonetheless a salient member of the same lexical class. Recent work has suggested that “none” is in fact a salient inferential alternative, at least in some models of implicature (?, ?). Thus, perhaps the strong correlation we observed between “some” and “none” is in fact due to the lack of a good semantics for “none” leading to failures. In other words, if children either don’t know or can’t process all the quantifiers in the lexical class, they may fail to reason appropriately about implicatures.

Several limitations in the current study provide directions for further work. First, we did not record children’s response time on the scalar implicature task. Perhaps a response latency measure might yield more information about the particular processing and inhibitory demands of scalar implicatures. Additionally, our quantifier knowledge measure (?, ?) may present many of the same pragmatic hurdles as our implicature task. For example, the syntactic construction of prompts in Give-Quantifier are very similar (e.g., “Can you put *some* of the oranges on the plate?” vs. “On the cover of my book, *some* of the pictures are oranges.”). It is also possible that, similarly to adults, children’s scalar implicature performance may also be affected by non-scalar alternatives, such as “one” or “two” (?, ?). Indeed, children show some evidence of comprehending quantifiers like “some” and “none” when presented in the same truth-value judgment task as exact number terms (?, ?). Future work could explore other more distinct measures of quantifier semantics, as well as the role of other alternatives involved in making scalar implicatures.

In sum, our work suggests that children can draw implicatures based on some lexical choices—such as in the case of ad-hoc implicatures—but they still struggle with quantifier-based scalar implicatures until relatively late. This trouble with quantifiers is not confined to making

scalar implicatures, but extends to children's knowledge of the quantifier "none" as well.

